Package 'BOLDx'

November 1, 2026

Type Package

Title Prediction and Time-to-Event Extensions of Ordered Lattice Dose Finding

Version 0.1.0

Author Wanru Guo [aut, cre], Gi-Ming Wang [ctb], Curtis Tatsuoka [ctb]

Maintainer Wanru Guo <wanru.guo@som.umaryland.edu>

Description Research simulation tools for prediction-assisted preventive intervention, delayed toxicity outcomes, and their combination in Bayesian ordered lattice dose finding. Includes complete-follow-up BOLD and BOIN comparators, a time-to-event BOIN comparator, configurable Beta priors, sequential child-specific mixture weights, and Monte Carlo operating characteristics. Extends the design of Wang and Tatsuoka (2026) <doi:10.1002/sim.70456>.

License GPL (>= 3)

Encoding UTF-8

Depends R (>= 4.1.0)

Imports Rcpp, stats

LinkingTo Rcpp

Suggests testthat (>= 3.0.0)

Config/testthat/edition 3

URL https://github.com/wguo3773/BOLDx

BugReports https://github.com/wguo3773/BOLDx/issues

RoxygenNote 7.3.3

NeedsCompilation yes

Contents

beta_mixture 2
bold_config 3
simulate_bold 4
simulate_scenarios 5
summarise_scenarios 6

Index **7**

`beta_mixture` *Inspect sequential weighted Beta components*

Description

Compute mixture weights, Beta component parameters, the mixture mean, and the overdose probability for prediction and partial-follow-up histories.

Usage

```
beta_mixture(
  a,
  b,
  q = numeric(),
  labels = rep(-1L, length(q)),
  sensitivity = 0.9,
  specificity = 0.98,
  target = 0.25
)
```

Arguments

a, b	Positive Beta parameters after resolved, nonlatent outcomes.
q	Remaining DLT-window fraction for each latent patient in [0,1]. Use 1 for an intervention patient without an observed DLT.
labels	Prediction labels: 1 positive, 0 negative, -1 uninformative.
sensitivity, specificity	Prediction operating characteristics.
target	Tail threshold for computing overdose probability.

Details

Components are weighted by sequential child-specific plug-in scores. Patient order is therefore part of the calculation. Known outcomes should first be incorporated into a and b. An observed intervention failure is a known DLT: increment a rather than branching.

Value

List containing components, mean, and overdose_probability.

Examples

```
beta_mixture(.75, 2.25, q = c(1, .5), labels = c(1L, -1L))
```

bold_config *Configure a BOLD extension simulation*

Description

The prior at dose j is Beta(prior_mean[j] * prior_pess[j], (1 - prior_mean[j]) * prior_pess[j]). True untreated risks are separate simulation inputs and are never supplied to dose-assignment calculations.

Usage

```
bold_config(
  target = 0.25,
  nmax = 21L,
  tau = 0.5,
  overdose_cutoff = 0.9,
  early_n = 12L,
  sensitivity = 0.9,
  specificity = 0.98,
  eligibility = 1,
  rho = 0.8,
  window = 3,
  accrual_rate = 2,
  arrival = c("uniform", "exponential"),
  timing = c("weibull", "uniform", "late"),
  prior_mean = NULL,
  prior_pess = 3,
  start_dose = 1L,
  pending_limit = 0.5,
  min_two = FALSE
)
```

Arguments

target	Target untreated DLT probability, strictly between 0 and 5/7.
nmax	Patient cap, a positive multiple of three.
tau	BOLD decision-tail target in (0,1); not the DLT target.
overdose_cutoff	Posterior overdose elimination cutoff in (0,1).
early_n	Stop when this many patients have received the current dose and the next recommendation is unchanged. Use Inf to disable.
sensitivity, specificity	Prediction sensitivity and specificity for latent untreated DLT status, each in [0,1].
eligibility	Probability a patient is eligible for prediction.
rho	Probability an intervention prevents an otherwise-occurring DLT.
window	DLT assessment window, in months.
accrual_rate	Mean number of arrivals per month.

arrival	Arrival-gap distribution, uniform or exponential.
timing	Generative conditional DLT-time distribution. The working follow-up model remains uniform for all choices.
prior_mean, prior_pess	Dose-specific vectors, or scalars recycled to all doses. NULL prior_mean uses target; PESS defaults to 3.
start_dose	Starting dose index, validated against each scenario.
pending_limit	Maximum pending fraction at the current dose before suspension; defaults to 0.5.
min_two	Require two completed outcomes before escalation.

Value

A validated configuration list of class bold_config.

Examples

```
cfg <- bold_config(nmax = 21, rho = 0.8)
cfg$target
```

---

simulate_bold *Simulate BOLD extensions and matched comparators*

Description

Evaluate dose-finding accuracy, enrollment, actual toxicity, allocation above the MTD, and calendar duration for BOLD extensions and their comparators.

Usage

```
simulate_bold(
  truth,
  method = c("prediction", "tite", "combined"),
  config = bold_config(),
  nsim = 100L,
  seed = 1L,
  mtd = NULL
)
```

Arguments

truth	Nondecreasing vector of 2 to 10 untreated DLT probabilities, strictly between zero and one.
method	One or more of prediction, tite, combined, bold, boin, tite_boin.
config	A configuration from bold_config().
nsim	Number of replicate trials per method.
seed	Nonnegative integer seed. The caller's random state is restored.
mtd	True optimal dose index for evaluation. NULL chooses the dose closest to target (lower index breaks ties). Set 0 for all-too-toxic; K+1 means all doses are below target and the highest dose is optimal. This value affects summaries only, never dose decisions.

Details

Fixed cohorts of three are used. Intervention failure is an observed DLT, not an extra latent branch. In combined mode, intervention follow-up resolves before the next dose decision. Mixture weights use sequential child-specific plug-in scores; the full calculation and assumptions are documented in the package's online Methods document.

Value

A bold_simulation list containing summary, trials, config, truth, mtd, and seed. Summary contains scenario-level point estimates. Optional Monte Carlo precision diagnostics are retained separately in mcse.

Examples

```
z <- simulate_bold(c(.10, .25, .45, .60),
  method = c("prediction", "bold", "boin"), nsim = 10, seed = 1)
z$summary
```

---

simulate_scenarios *Run a panel of dose-toxicity scenarios*

Description

Results include scenario-level estimates and mean/SD summaries across random scenarios within each method and MTD category. Use explicit MTD labels for all-too-toxic and below-range cases.

Usage

```
simulate_scenarios(
  scenarios,
  mtd = NULL,
  method = c("prediction", "tite", "combined"),
  config = bold_config(),
  nsim = 100,
  seed = 1
)
```

Arguments

scenarios	Numeric matrix with one nondecreasing risk vector per row.
mtd	Optional integer vector giving each scenario's optimal dose.
method, config, nsim, seed	As in simulate_bold(). Seeds increment by scenario index; methods within a scenario share the same seed.

Value

List with summary and trials data frames, each containing scenario, and aggregate containing means and between-scenario standard deviations.

Examples

```
x <- rbind(c(.1,.25,.45,.6),c(.05,.1,.25,.45))
simulate_scenarios(x, mtd=c(2,3), method="prediction", nsim=5)$summary
```

`summarise_scenarios` *Summarise operating characteristics across random scenarios*

Description

Compute the equally weighted mean and sample standard deviation of scenario-level estimates, each obtained from its simulated trial replicates.

Usage

```
summarise_scenarios(x, group_by = c("method", "mtd"))
```

Arguments

`x`	A data frame of scenario-level summaries, or the result returned by simulate_scenarios().
`group_by`	Columns defining the reporting groups; must include method. Additional columns, such as delta, can be attached to x before summarising.

Details

The SD measures between-scenario dispersion, not a standard error. For 50 scenarios with 100 replicates each, the SD is calculated from the 50 scenario-level estimates, not from the 5,000 individual trials. It includes both scenario heterogeneity and finite-replicate simulation noise. Group by MTD category, upper toxicity gap, and design settings where appropriate. A group with only one scenario has undefined SD, reported as NA. Above-MTD percentages are first patient-weighted within each scenario, then equally weighted across scenarios. Missing duration for complete-follow-up methods is retained as NA.

Value

A data frame of point-estimate means and matching *_sd columns, with n_scenarios and the minimum/maximum replicates per scenario.

Calculation

For J scenario-level estimates Q_j, the reported values are

$$\overline{Q} = J^{-1}\sum_{j=1}^{J} Q_j, \qquad SD(Q) = \sqrt{\frac{\sum_{j=1}^{J}(Q_j - \overline{Q})^2}{J-1}}.$$

Examples

```
panel <- simulate_scenarios(rbind(c(.1,.25,.45), c(.15,.25,.4)),
  mtd=c(2,2), method="prediction", nsim=5)
panel$aggregate
```

Index

beta_mixture, 2
bold_config, 3

simulate_bold, 4
simulate_scenarios, 5
summarise_scenarios, 6